

Installing Hadoop 3.2.1 Single node cluster on Windows 10



Hadi Fadlallah

Follow

Apr 18 · 7 min read

While working on a project two years ago, I wrote a step-by-step guide to install Hadoop 3.1.0 on Ubuntu 16.04 operating system. Since we are currently working on a new project where we need to install a Hadoop cluster on Windows 10, I decided to write a guide for this process.

1. Prerequisites

First, we need to make sure that the following prerequisites are installed:

1. Java 8 runtime environment (JRE): Hadoop 3 requires a Java 8 installation. I prefer using the offline installer.

2. Java 8 development Kit (JDK).

3. To unzip downloaded Hadoop binaries, we should install 7zip.

4. I will create a folder “E:\hadoop-env” on my local machine to store downloaded files.

2. Download Hadoop binaries

The first step is to download Hadoop binaries from the official website. The binary package size is about 342 MB.



Figure 1 — Hadoop binaries download link

After finishing the file download, we should unpack the package using 7zip in two steps. First, we should extract the hadoop-3.2.1.tar.gz library, and

then, we should unpack the extracted tar file:

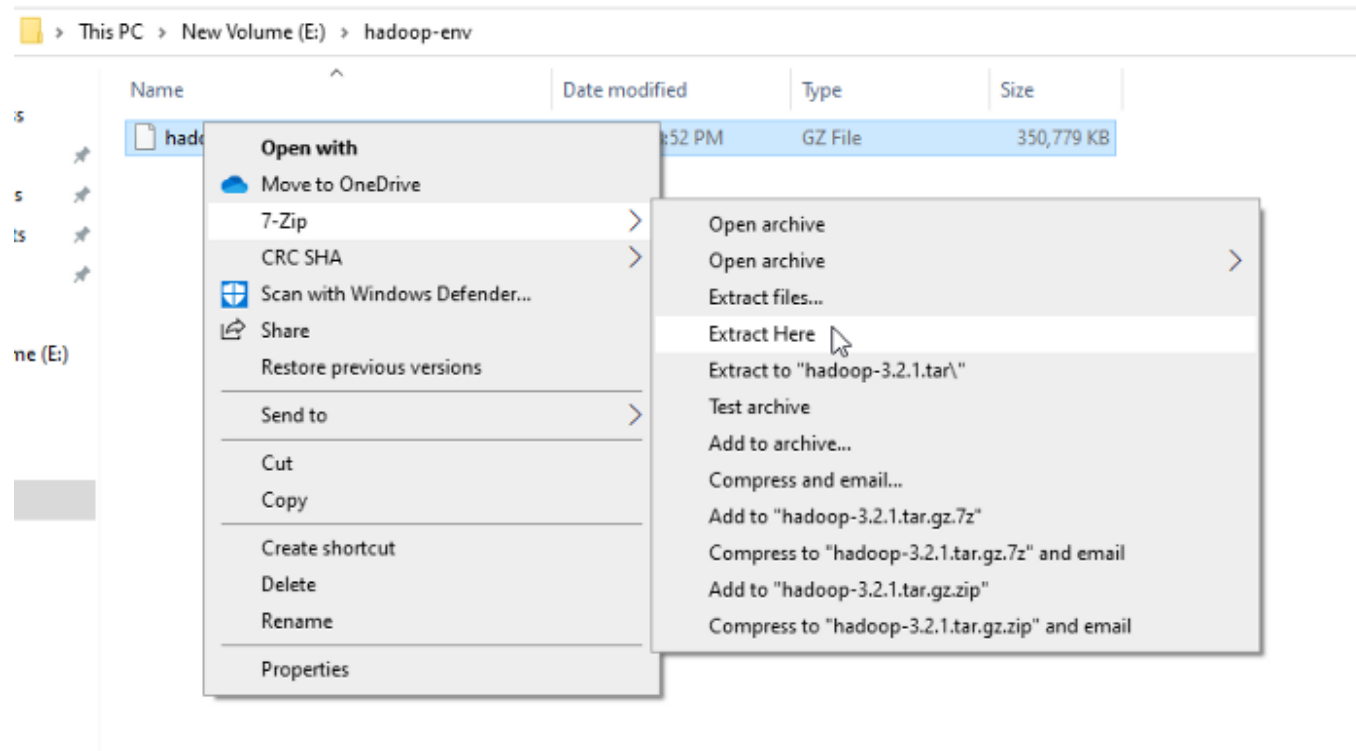


Figure 2 — Extracting hadoop-3.2.1.tar.gz package using 7zip

Name	Date modified	Type	Size
hadoop-3.2.1.tar	9/10/2019 8:11 PM	TAR File	893,250 KB
hadoop-3.2.1.tar.gz	4/15/2020 8:52 PM	GZ File	350,779 KB

Figure 3 — Extracted hadoop-3.2.1.tar file



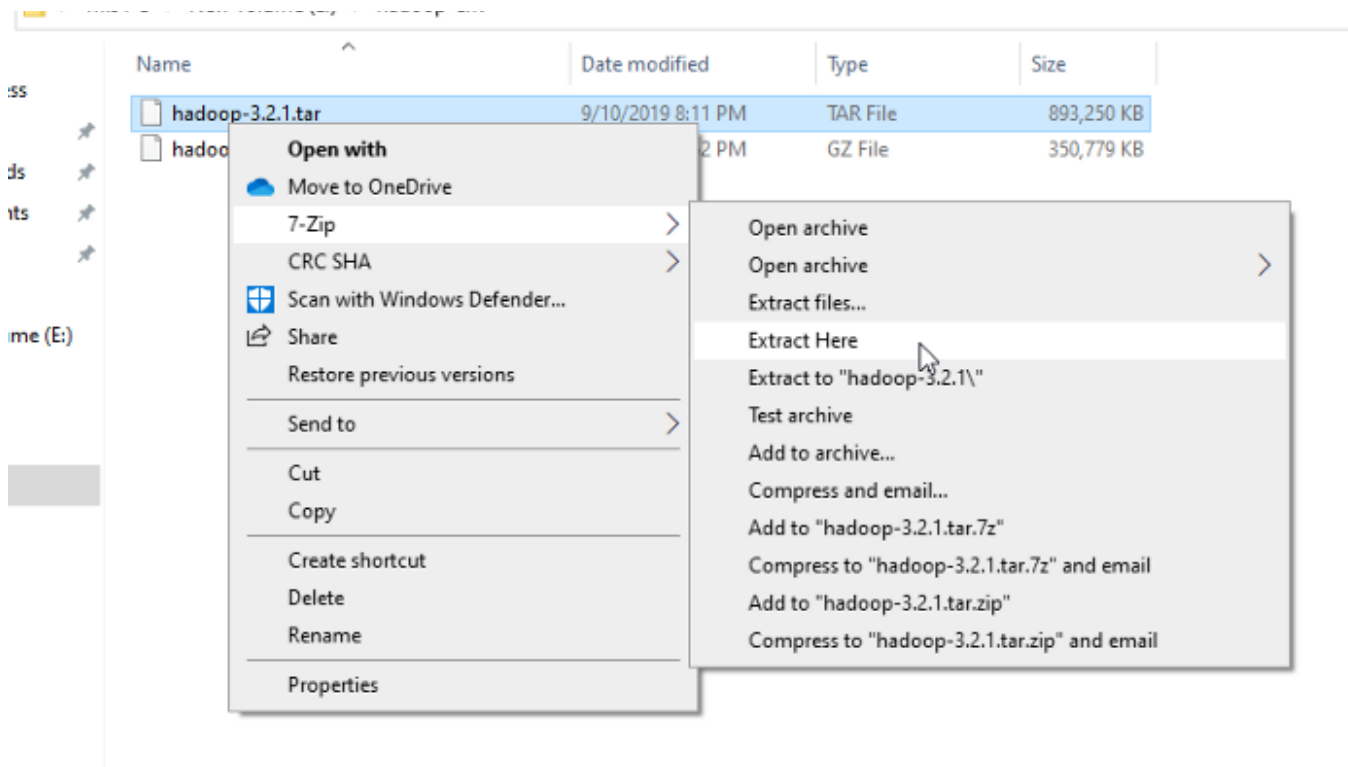
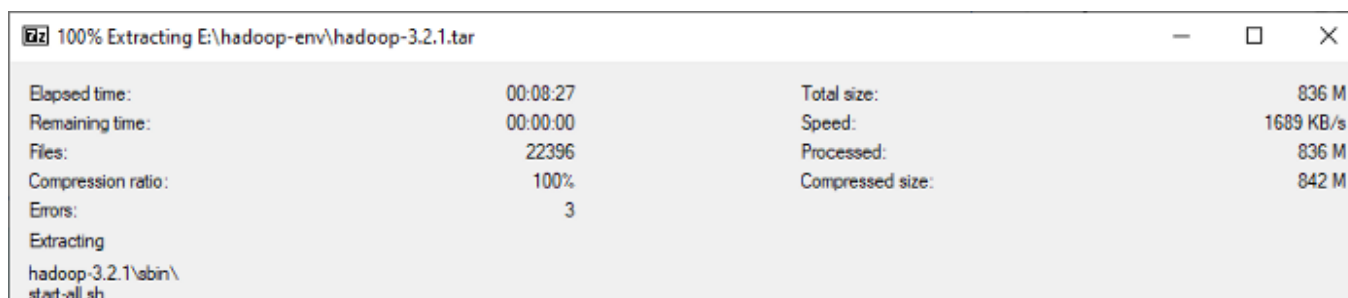


Figure 4 — Extracting the hadoop-3.2.1.tar file

The tar file extraction may take some minutes to finish. In the end, you may see some warnings about symbolic link creation. Just ignore these warnings since they are not related to windows.



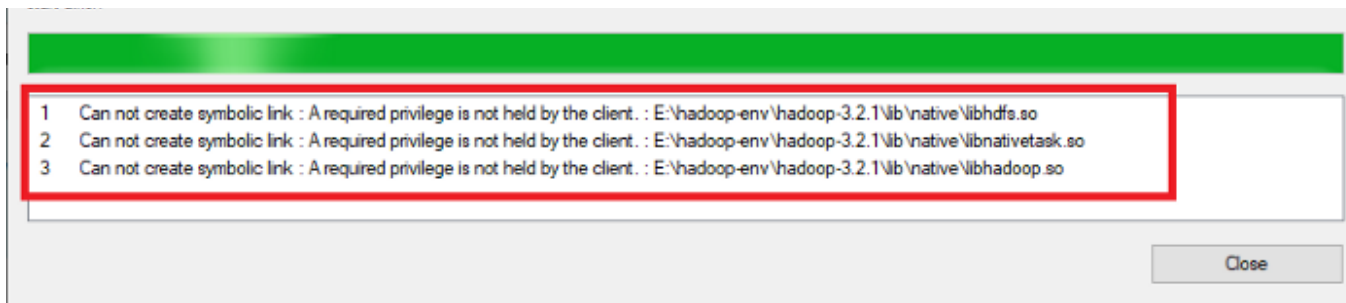


Figure 5 — Symbolic link warnings

After unpacking the package, we should add the Hadoop native IO libraries, which can be found in the following GitHub repository:

<https://github.com/cdarlint/winutils>.

Since we are installing Hadoop 3.2.1, we should download the files located in <https://github.com/cdarlint/winutils/tree/master/hadoop-3.2.1/bin> and copy them into the “hadoop-3.2.1\bin” directory.

3. Setting up environment variables

After installing Hadoop and its prerequisites, we should configure the environment variables to define Hadoop and Java default paths.

To edit environment variables, go to Control Panel > System and Security > System (or right-click > properties on My Computer icon) and click on the “Advanced system settings” link.

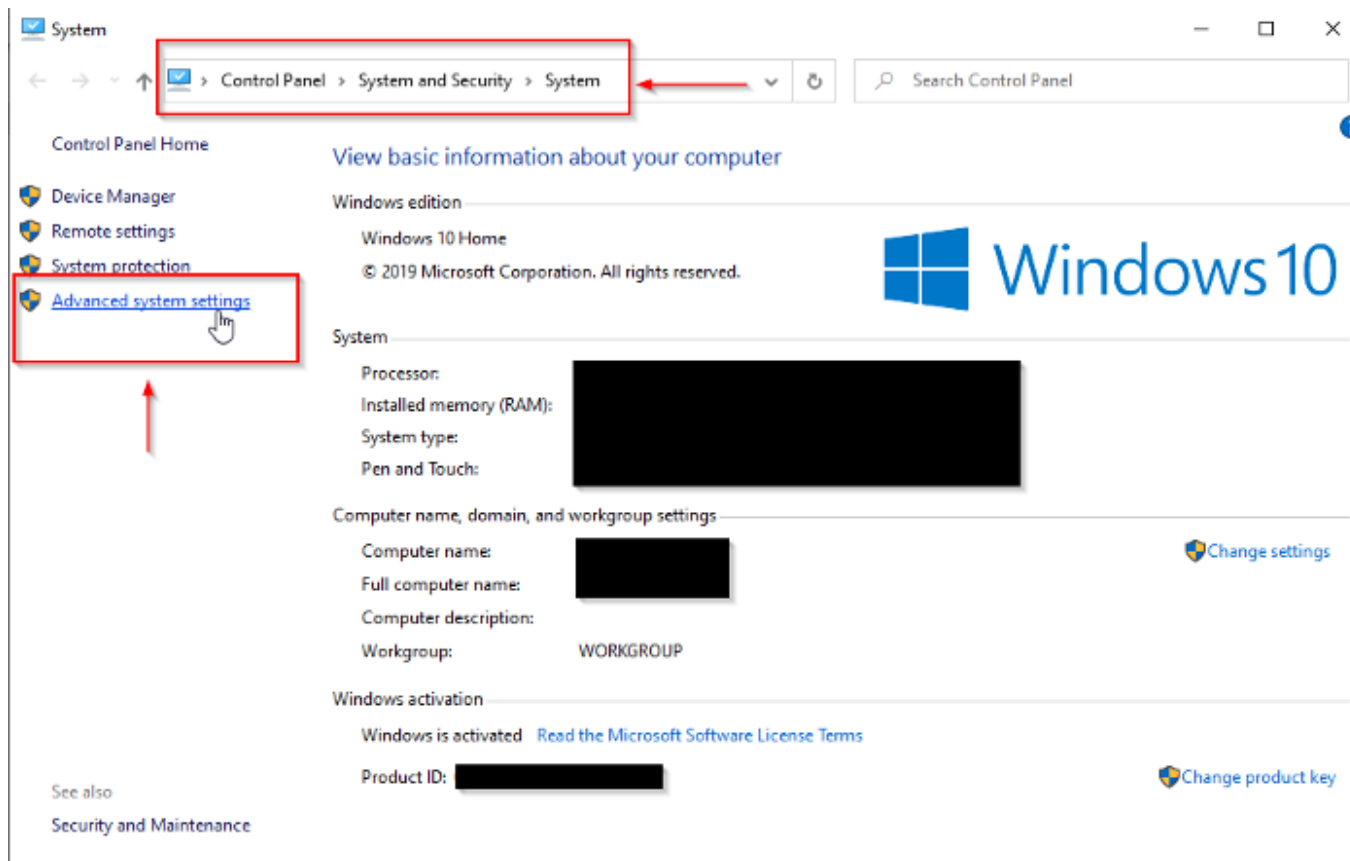



Figure 6 — Opening advanced system settings

When the “Advanced system settings” dialog appears, go to the “Advanced” tab and click on the “Environment variables” button located on the bottom of the dialog.



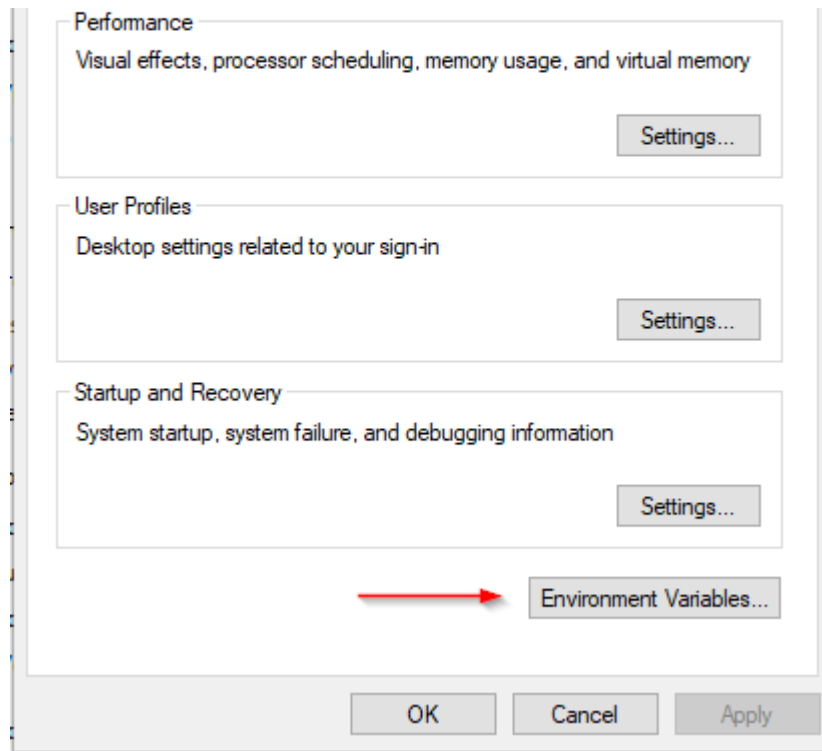


Figure 7 — Advanced system settings dialog

In the “Environment Variables” dialog, press the “New” button to add a new variable.

Note: In this guide, we will add user variables since we are configuring Hadoop for a single user. If you are looking to configure Hadoop for multiple users, you can define System variables instead.

There are two variables to define:

1. JAVA_HOME: JDK installation folder path

2. HADOOP_HOME: Hadoop installation folder path

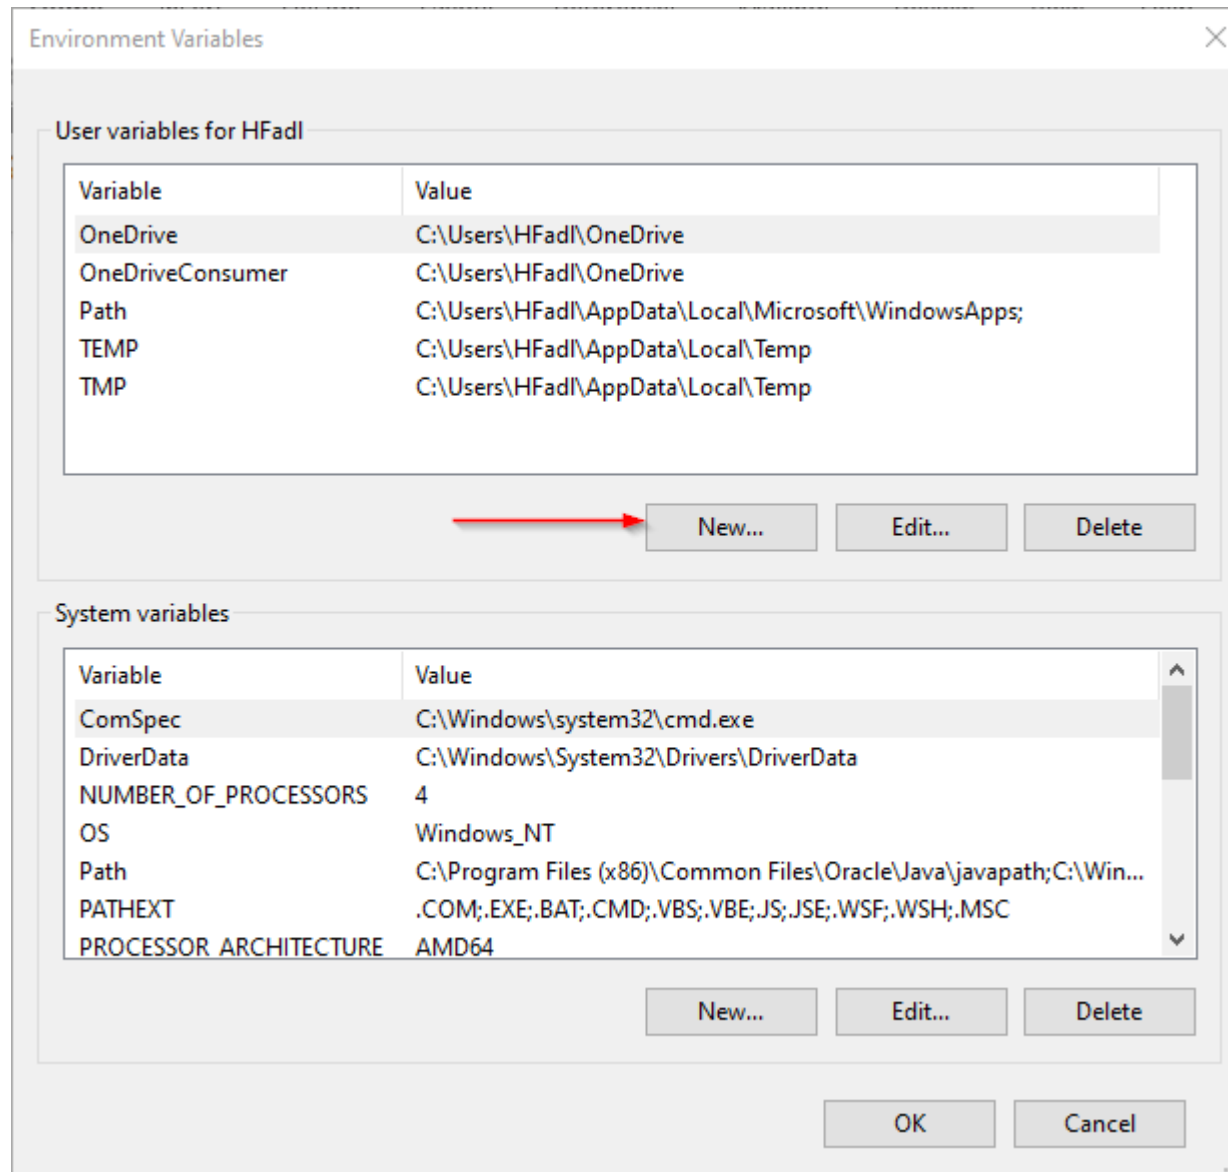


Figure 8 — Adding JAVA_HOME variable




Figure 9 — Adding HADOOP_HOME variable

Now, we should edit the PATH variable to add the Java and Hadoop binaries paths as shown in the following screenshots.





Figure 10 — Editing the PATH variable



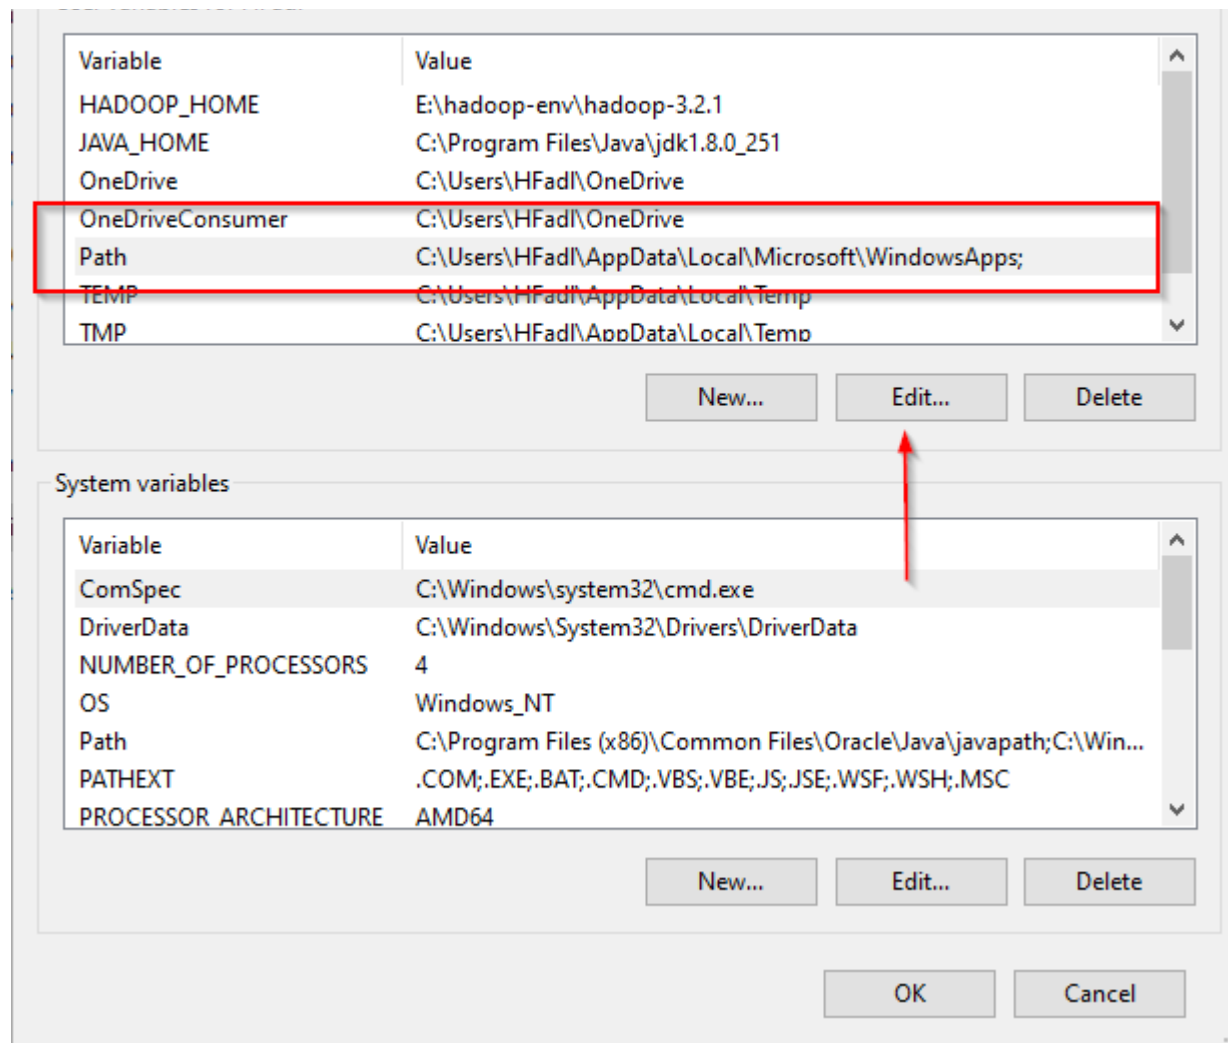
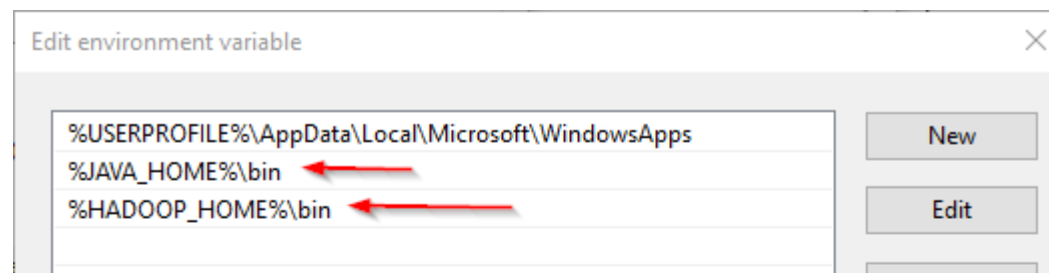


Figure 11 — Editing PATH variable






Figure 12— Adding new paths to the PATH variable

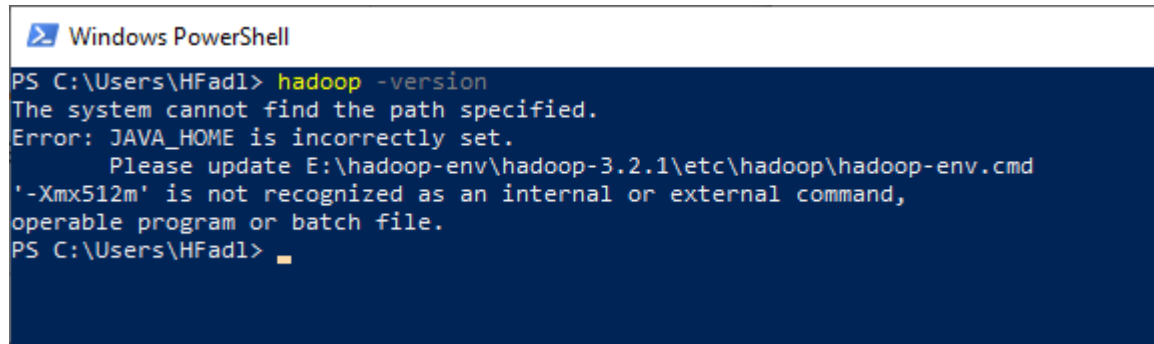
3.1. JAVA_HOME is incorrectly set error

Now, let's open PowerShell and try to run the following command:

```
hadoop -version
```

In this example, since the JAVA_HOME path contains spaces, I received the following error:

JAVA_HOME is incorrectly set

A screenshot of a Windows PowerShell window. The title bar says "Windows PowerShell". The command prompt shows the user at C:\Users\HFadl> typing 'hadoop -version'. The output is an error: "The system cannot find the path specified. Error: JAVA_HOME is incorrectly set. Please update E:\hadoop-env\hadoop-3.2.1\etc\hadoop\hadoop-env.cmd '-Xmx512m' is not recognized as an internal or external command, operable program or batch file." The prompt returns to C:\Users\HFadl>.

```
PS C:\Users\HFadl> hadoop -version
The system cannot find the path specified.
Error: JAVA_HOME is incorrectly set.
    Please update E:\hadoop-env\hadoop-3.2.1\etc\hadoop\hadoop-env.cmd
'-Xmx512m' is not recognized as an internal or external command,
operable program or batch file.
PS C:\Users\HFadl>
```

Figure 13 — JAVA_HOME error

To solve this issue, we should use the windows 8.3 path instead. As an example:

- Use “Progra~1” instead of “Program Files”
- Use “Progra~2” instead of “Program Files(x86)”

After replacing “Program Files” with “Progra~1”, we closed and reopened PowerShell and tried the same command. As shown in the screenshot below, it runs without errors.

A partial screenshot of a Windows PowerShell window. The title bar is visible, showing the PowerShell icon and the text "Windows PowerShell".

```
Windows PowerShell
```

```
Copyright (C) Microsoft Corporation. All rights reserved.  
  
Try the new cross-platform PowerShell https://aka.ms/pscore6  
  
PS C:\Users\HFadl> hadoop -version  
java version "1.8.0_251"  
Java(TM) SE Runtime Environment (build 1.8.0_251-b08)  
Java HotSpot(TM) 64-Bit Server VM (build 25.251-b08, mixed mode)  
PS C:\Users\HFadl>
```

Figure 14 — `hadoop -version` command executed successfully

4. Configuring Hadoop cluster

There are four files we should alter to configure Hadoop cluster:

1. %HADOOP_HOME%\etc\hadoop\hdfs-site.xml
2. %HADOOP_HOME%\etc\hadoop\core-site.xml
3. %HADOOP_HOME%\etc\hadoop\mapred-site.xml
4. %HADOOP_HOME%\etc\hadoop\yarn-site.xml

4.1. HDFS site configuration

As we know, Hadoop is built using a master-slave paradigm. Before altering the HDFS configuration file, we should create a directory to store all master node (name node) data and another one to store data (data node). In this example, we created the following directories:

- E:\hadoop-env\hadoop-3.2.1\data\dfs\namenode
- E:\hadoop-env\hadoop-3.2.1\data\dfs\datanode

Now, let's open "hdfs-site.xml" file located in
"%HADOOP_HOME%\etc\hadoop" directory, and we should add the
following properties within the <configuration> </configuration>
element:

```
<property>
```

```
<name>dfs.replication</name>
```

```
<value>1</value>
```

```
</property>
```

```
<property>
```

```
<name>dfs.namenode.name.dir</name>
```

```
<value>file:///E:/hadoop-env/hadoop-3.2.1/data/dfs/namenode</value>
```

```
</property>
```

```
<property>
```

```
<name>dfs.datanode.data.dir</name>

<value>file:///E:/hadoop-env/hadoop-3.2.1/data/dfs/datanode</value>

</property>
```

Note that we have set the replication factor to 1 since we are creating a single node cluster.

4.2. Core site configuration

Now, we should configure the name node URL adding the following XML code into the `<configuration>` `</configuration>` element within “core-site.xml”:

```
<property>

<name>fs.default.name</name>

<value>hdfs://localhost:9820</value>

</property>
```

4.3. Map Reduce site configuration

Now, we should add the following XML code into the `<configuration>`
`</configuration>` element within “mapred-site.xml”:

```
<property>

<name>mapreduce.framework.name</name>

<value>yarn</value>

<description>MapReduce framework name</description>

</property>
```

4.4. Yarn site configuration

Now, we should add the following XML code into the `<configuration>`
`</configuration>` element within “yarn-site.xml”:

```
<property>

<name>yarn.nodemanager.aux-services</name>

<value>mapreduce_shuffle</value>

<description>Yarn Node Manager Aux Service</description>
```



```
</property>
```

5. Formatting Name node

After finishing the configuration, let's try to format the name node using the following command:

```
hdfs namenode -format
```

Due to a bug in the Hadoop 3.2.1 release, you will receive the following error:

```
2020-04-17 22:04:01,503 ERROR namenode.NameNode: Failed to start  
namenode.
```

```
java.lang.UnsupportedOperationException
```

```
at java.nio.file.Files.setPosixFilePermissions(Files.java:2044)
```

```
at  
org.apache.hadoop.hdfs.server.common.Storage$StorageDirectory.clearDi  
rectory(Storage.java:452)
```

```
at
org.apache.hadoop.hdfs.server.namenode.NNStorage.format (NNStorage.java:591)

at
org.apache.hadoop.hdfs.server.namenode.NNStorage.format (NNStorage.java:613)

at
org.apache.hadoop.hdfs.server.namenode.FSImage.format (FSImage.java:188)

at
org.apache.hadoop.hdfs.server.namenode.NameNode.format (NameNode.java:1206)

at
org.apache.hadoop.hdfs.server.namenode.NameNode.createNameNode (NameNode.java:1649)

at
org.apache.hadoop.hdfs.server.namenode.NameNode.main (NameNode.java:1759)

2020-04-17 22:04:01,511 INFO util.ExitUtil: Exiting with status 1:
java.lang.UnsupportedOperationException

2020-04-17 22:04:01,518 INFO namenode.NameNode: SHUTDOWN_MSG:
```

This issue will be solved within the next release. For now, you can fix it temporarily using the following steps ([reference](#)):

1. Download hadoop-hdfs-3.2.1.jar file from the [following link](#).
2. Rename the file name hadoop-hdfs-3.2.1.jar to hadoop-hdfs-3.2.1.bak in folder %HADOOP_HOME%\share\hadoop\hdfs
3. Copy the downloaded hadoop-hdfs-3.2.1.jar to folder %HADOOP_HOME%\share\hadoop\hdfs

Now, if we try to re-execute the format command (Run the command prompt or PowerShell as administrator), you need to approve file system format.

```
2020-04-17 22:02:58,422 INFO util.GSet: Computing capacity for map NameNodeRetryCache
2020-04-17 22:02:58,423 INFO util.GSet: VM type = 64-bit
2020-04-17 22:02:58,424 INFO util.GSet: 0.0299999999329447746% max memory 889 MB = 273.1 KB
2020-04-17 22:02:58,425 INFO util.GSet: capacity = 2^15 = 32768 entries
Re-format filesystem in Storage Directory root= E:\hadoop-env\hadoop-3.2.1\data\dfs\namenode; location= null ? (Y or N)
y
```

Figure 15 — File system format approval

And the command is executed successfully:

```
2020-04-17 22:14:17,206 INFO namenode.FSImage: Allocated new BlockPoolId: BP-2032026115-192.168.1.105-1587150857190
2020-04-17 22:14:17,207 INFO common.Storage: Will remove files: []
2020-04-17 22:14:17,275 INFO common.Storage: Storage directory E:\hadoop-env\hadoop-3.2.1\data\dfs\namenode has been successfully formatted.
2020-04-17 22:14:17,331 INFO namenode.FSImageFormatProtobuf: Saving image file E:\hadoop-env\hadoop-3.2.1\data\dfs\namenode\current\fsimage.ckpt_000000000000000000 using no compression
2020-04-17 22:14:17,531 INFO namenode.FSImageFormatProtobuf: Image file E:\hadoop-env\hadoop-3.2.1\data\dfs\namenode\current\fsimage.ckpt_000000000000000000 of size 400 bytes saved in 0 seconds .
2020-04-17 22:14:17,555 INFO namenode.NNStorageRetentionManager: Going to retain 1 images with txid >= 0
2020-04-17 22:14:17,580 INFO namenode.FSImage: FSImageSaver clean checkpoint: txid=0 when meet shutdown.
2020-04-17 22:14:17,580 INFO namenode.NameNode: SHUTDOWN_MSG:
*****
```

```
SHUTDOWN_MSG: Shutting down NameNode at [REDACTED]
*****
PS C:\Windows\system32>
```

Figure 16 — Command executed successfully

6. Starting Hadoop services

Now, we will open PowerShell, and navigate to “%HADOOP_HOME%\sbin” directory. Then we will run the following command to start the Hadoop nodes:

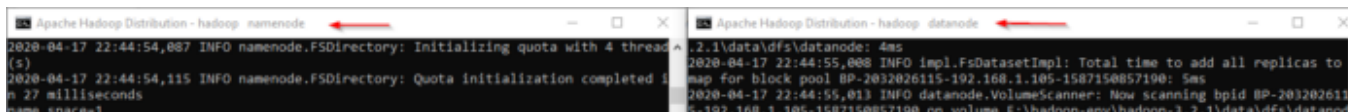
```
.\start-dfs.cmd
```



```
Administrator: Windows PowerShell
PS E:\hadoop-env\hadoop-3.2.1\sbin> .\start-dfs.cmd
PS E:\hadoop-env\hadoop-3.2.1\sbin>
```

Figure 17 — Starting Hadoop nodes

Two command prompt windows will open (one for the name node and one for the data node) as follows:



```
Apache Hadoop Distribution - hadoop namenode
2020-04-17 22:44:54,087 INFO namenode.FSDirectory: Initializing quota with 4 threads
2020-04-17 22:44:54,115 INFO namenode.FSDirectory: Quota initialization completed in 27 milliseconds
name space=1

Apache Hadoop Distribution - hadoop datanode
2020-04-17 22:44:55,088 INFO impl.FsDatasetImpl: Total time to add all replicas to map for block pool BP-2032026115-192.168.1.105-1587150857190: 5ms
2020-04-17 22:44:55,013 INFO datanode.VolumeScanner: Now scanning bpid BP-2032026115-192.168.1.105-1587150857190 on volume f:\hadoop-env\hadoop-3.2.1\data\dfs\data...
```

```
storage space=0
storage types=RAM_DISK=0, SSD=0, DISK=0, ARCHIVE=0, PROVIDED=0
2020-04-17 22:44:54,151 INFO blockmanagement.CacheReplicationMonitor: Starting CacheReplicationMonitor with interval 30000 milliseconds
2020-04-17 22:44:55,147 INFO hdfs.StateChange: BLOCK* registerDatanode: from DatanodeRegistration(127.0.0.1:9866, datanodeUid=94e235fa-78fd-413e-95e5-5d84dc62bc9f, infoPort=9864, infoSecurePort=0, ipcPort=9867, storageInfo=lv=-57;cid=CID-11d9a863-fc7b-4208-b192-2e4b6bf8736f;nsid=668255427;c=1587150857190) storage 94e235fa-78fd-413e-95e5-5d84dc62bc9f
2020-04-17 22:44:55,152 INFO net.NetworkTopology: Adding a new node: /default-rack/127.0.0.1:9866
2020-04-17 22:44:55,153 INFO blockmanagement.BlockReportHandler: Registered Datanode 94e235fa-78fd-413e-95e5-5d84dc62bc9f (127.0.0.1:9866).
2020-04-17 22:44:55,353 INFO blockmanagement.DatanodeDescriptor: Adding new storage ID DS-de0ef9eb-f03b-40b4-8bdd-dd36b16ee068 for DN 127.0.0.1:9866
2020-04-17 22:44:55,473 INFO BlockStateChange: BLOCK* processReport @x6718670216286c90: Processing first storage report for DS-de0ef9eb-f03b-40b4-8bdd-dd36b16ee068 from datanode 94e235fa-78fd-413e-95e5-5d84dc62bc9f
2020-04-17 22:44:55,478 INFO BlockStateChange: BLOCK* processReport @x6718670216286c90: from storage DS-de0ef9eb-f03b-40b4-8bdd-dd36b16ee068 node DatanodeRegistration(127.0.0.1:9866, datanodeUid=94e235fa-78fd-413e-95e5-5d84dc62bc9f, infoPort=9864, infoSecurePort=0, ipcPort=9867, storageInfo=lv=-57;cid=CID-11d9a863-fc7b-4208-b192-2e4b6bf8736f;nsid=668255427;c=1587150857190), blocks: 0, hasStaleStorage: false, processing time: 5 msec, invalidatedBlocks: 0
2020-04-17 22:44:55,016 INFO datanode.VolumeScanner: VolumeScanner(E:\hadoop-env\hadoop-3.2.1\data\dfs\datanode, DS-de0ef9eb-f03b-40b4-8bdd-dd36b16ee068): finished scanning block pool BP-2032026115-192.168.1.105-1587150857190
2020-04-17 22:44:55,059 INFO datanode.DirectoryScanner: Periodic Directory Tree Verification scan starting at 4/18/20 1:51 AM with interval of 21600000ms
2020-04-17 22:44:55,069 INFO datanode.VolumeScanner: VolumeScanner(E:\hadoop-env\hadoop-3.2.1\data\dfs\datanode, DS-de0ef9eb-f03b-40b4-8bdd-dd36b16ee068): no suitable block pools found to scan. Waiting 1814399943 ms.
2020-04-17 22:44:55,075 INFO datanode.DataNode: Block pool BP-2032026115-192.168.1.105-1587150857190 (Datanode Uid 94e235fa-78fd-413e-95e5-5d84dc62bc9f) service to localhost/127.0.0.1:9820 beginning handshake with NN
2020-04-17 22:44:55,182 INFO datanode.DataNode: Block pool Block pool BP-2032026115-192.168.1.105-1587150857190 (Datanode Uid 94e235fa-78fd-413e-95e5-5d84dc62bc9f) service to localhost/127.0.0.1:9820 successfully registered with NN
2020-04-17 22:44:55,183 INFO datanode.DataNode: For namenode localhost/127.0.0.1:9820 using BLOCKREPORT_INTERVAL of 21600000msec CACHEREPORT_INTERVAL of 10000msec Initial delay: 0msec; heartbeatInterval=3000
2020-04-17 22:44:55,555 INFO datanode.DataNode: Successfully sent block report @x6718670216286c90, containing 1 storage report(s), of which we sent 1. The reports had 0 total blocks and used 1 RPC(s). This took 12 msec to generate and 129 msec for RPC and NN processing. Got back one command: FinalizeCommand/5.
2020-04-17 22:44:55,556 INFO datanode.DataNode: Got finalize command for block pool BP-2032026115-192.168.1.105-1587150857190
```

Figure 18 — Hadoop nodes command prompt windows

Next, we must start the Hadoop Yarn service using the following command:

```
./start-yarn.cmd
```

```
PS E:\hadoop-env\hadoop-3.2.1\sbin> .\start-yarn.cmd
starting yarn daemons
PS E:\hadoop-env\hadoop-3.2.1\sbin>
```

Figure 19 — Starting Hadoop Yarn services

Two command prompt windows will open (one for the resource manager and one for the node manager) as follows:



```

2020-04-17 22:47:05,659 INFO store.AbstractFSNodeStore: Finished create editing file
at:file:/tmp/hadoop-yarn-HFadl/node-attribute/nodeattribute.editlog
2020-04-17 22:47:05,695 INFO event.AsyncDispatcher: Registering class org.apache.ha
doo
p.yarn.server.resourcemanager.node.labels.NodeAttributesStoreEventType for class
org.apache.hadoop.yarn.server.resourcemanager.node.labels.NodeAttributesManagerImpl
ForwardingEventHandler
2020-04-17 22:47:05,782 INFO placement.MultiNodeSortingManager: Starting NodeSortin
gService-MultiNodeSortingManager
2020-04-17 22:47:05,749 INFO ipc.CallQueueManager: Using callQueue: class java.util
.concurrent.LinkedBlockingQueue, queueCapacity: 5000, scheduler: class org.apache.h
ad
oop.ipc.DefaultRpcScheduler, ipcBackoff: false.
2020-04-17 22:47:05,765 INFO pb.RpcServerFactoryPBImpl: Adding protocol org.apache.
hadoop.yarn.server.api.ResourceTrackerPB to the server
2020-04-17 22:47:05,799 INFO ipc.Server: Starting Socket Reader #1 for port 0031
2020-04-17 22:47:05,825 INFO ipc.Server: IPC Server listener on 0031: starting
2020-04-17 22:47:05,828 INFO ipc.Server: IPC Server Responder: starting
2020-04-17 22:47:05,847 INFO util.JvmPauseMonitor: Starting JVM pause monitor
2020-04-17 22:47:05,853 INFO ipc.CallQueueManager: Using callQueue: class java.util
.concurrent.LinkedBlockingQueue, queueCapacity: 5000, scheduler: class org.apache.h
ad
oop.ipc.DefaultRpcScheduler, ipcBackoff: false.
2020-04-17 22:47:05,929 INFO ipc.Server: Starting Socket Reader #1 for port 0030
2020-04-17 22:47:05,956 INFO pb.RpcServerFactoryPBImpl: Adding protocol org.apache.
hadoop.yarn.api.ApplicationMasterProtocolPB to the server
2020-04-17 22:47:06,007 INFO ipc.Server: IPC Server listener on 0030: starting
2020-04-17 22:47:06,008 INFO ipc.Server: IPC Server Responder: starting
2020-04-17 22:47:06,263 INFO ipc.CallQueueManager: Using callQueue: class java.util
.concurrent.LinkedBlockingQueue, queueCapacity: 5000, scheduler: class org.apache.h
ad
oop.ipc.DefaultRpcScheduler, ipcBackoff: false.
2020-04-17 22:47:06,288 INFO ipc.Server: Starting Socket Reader #1 for port 0032
2020-04-17 22:47:06,298 INFO pb.RpcServerFactoryPBImpl: Adding protocol org.apache.
hadoop.yarn.api.ApplicationClientProtocolPB to the server
2020-04-17 22:47:06,320 INFO resourcemanager.ResourceManager: Transitioned to activ
e state
2020-04-17 22:47:06,331 INFO ipc.Server: IPC Server Responder: starting
2020-04-17 22:47:06,333 INFO ipc.Server: IPC Server listener on 0032: starting
2020-04-17 22:47:06,961 INFO resourcemanager.ResourceTrackerService: NodeManager fr
om node DESKTOP-SSVATPQ(cwPort: 57849 httpPort: 8042) registered with capability: <
memory:8192, vCores:8>, assigned nodeId DESKTOP-SSVATPQ:57849
2020-04-17 22:47:06,978 INFO rmnode.RMNodeImpl: DESKTOP-SSVATPQ:57849 Node Transiti
oned from NEW to RUNNING
2020-04-17 22:47:07,016 INFO capacity.CapacityScheduler: Added node DESKTOP-SSVATP
Q:57849 clusterResource: <memory:8192, vCores:8>
Apr 17, 2020 10:47:04 PM com.sun.jersey.guice.spi.container.GuiceComponentProvider
factory register
INFO: Registering org.apache.hadoop.yarn.server.nodemanager.webapp.JAXContextResol
ver as a provider class
Apr 17, 2020 10:47:04 PM com.sun.jersey.server.impl.application.WebApplicationImpl
Initiate
INFO: Initiating Jersey application, version 'Jersey: 1.19 02/11/2015 03:25 AM'
Apr 17, 2020 10:47:04 PM com.sun.jersey.guice.spi.container.GuiceComponentProvid
er factory getComponentProvider
INFO: Binding org.apache.hadoop.yarn.server.nodemanager.webapp.JAXContextResol
ver to GuiceManagedComponentProvider with the scope "Singleton"
Apr 17, 2020 10:47:04 PM com.sun.jersey.guice.spi.container.GuiceComponentProvid
er factory getComponentProvider
INFO: Binding org.apache.hadoop.yarn.webapp.GenericExceptionHandler to GuiceMan
agedComponentProvider with the scope "Singleton"
Apr 17, 2020 10:47:06 PM com.sun.jersey.guice.spi.container.GuiceComponentProvid
er factory getComponentProvider
INFO: Binding org.apache.hadoop.yarn.server.nodemanager.webapp.NMWebServices to
GuiceManagedComponentProvider with the scope "Singleton"
2020-04-17 22:47:06,268 INFO handler.ContextHandler: Started o.e.j.w.WebAppContext@
716b99a9[/,file:///C:/Users/HFadl/AppData/Local/Temp/jetty-0.0.0-0-8042-node-_-any-
135495797710539143.dlr/webapp/,AVAILABLE!]{/node}
2020-04-17 22:47:06,284 INFO server.AbstractConnector: Started ServerConnector@06f
5a8(HTTP/1.1,[http/1.1]){0.0.0.0:8042}
2020-04-17 22:47:06,285 INFO server.Server: Started @13636ms
2020-04-17 22:47:06,286 INFO webapp.Webapps: Web app node started at 8042
2020-04-17 22:47:06,289 INFO nodemanager.NodeStatusUpdaterImpl: Node ID assigned is
: DESKTOP-SSVATPQ:57849
2020-04-17 22:47:06,319 INFO client.NMProxy: Connecting to ResourceManager at /0.0.
0.0:8031
2020-04-17 22:47:06,327 INFO util.JvmPauseMonitor: Starting JVM pause monitor
2020-04-17 22:47:06,448 INFO nodemanager.NodeStatusUpdaterImpl: Sending out 0 NM co
ntainer statuses: []
2020-04-17 22:47:06,476 INFO nodemanager.NodeStatusUpdaterImpl: Registering with RM
using containers: []
2020-04-17 22:47:06,999 INFO security.NMContainerTokenSecretManager: Rolling master
-key for container-tokens, got key with id -611578012
2020-04-17 22:47:07,001 INFO security.NMTokenSecretManagerImpl: Rolling master-key
for container-tokens, got key with id -977682056
2020-04-17 22:47:07,003 INFO nodemanager.NodeStatusUpdaterImpl: Registered with Res
ourceManager as DESKTOP-SSVATPQ:57849 with total resource of <memory:8192, vCores:8>

```

Figure 20— Node manager and Resource manager command prompt windows

To make sure that all services started successfully, we can run the following command:

```
jps
```

It should display the following services:

```
14560 DataNode
4960 ResourceManager
5936 NameNode
768 NodeManager
14636 Jps
```

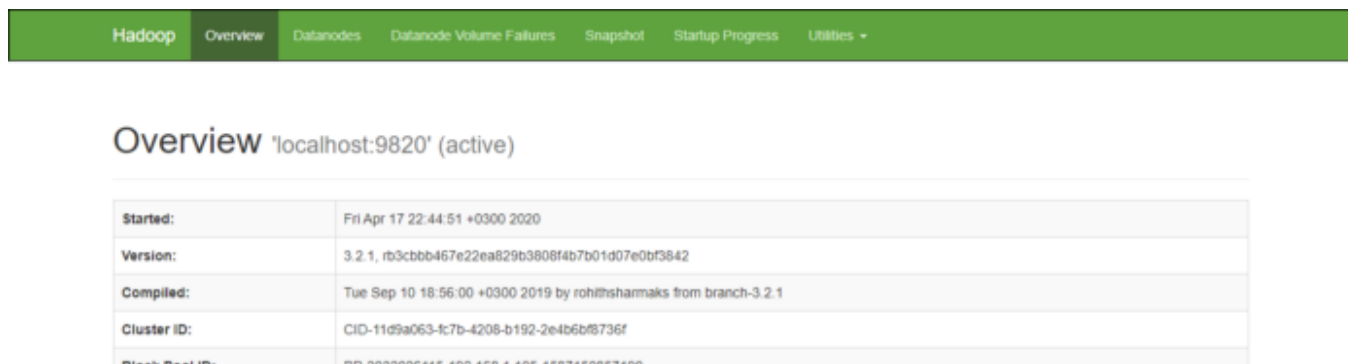
```
PS E:\hadoop-env\hadoop-3.2.1\sbin> jps
14560 DataNode
4960 ResourceManager
5936 NameNode
768 NodeManager
14636 Jps
PS E:\hadoop-env\hadoop-3.2.1\sbin>
```

Figure 21 — Executing jps command

7. Hadoop Web UI

There are three web user interfaces to be used:

- Name node web page: <http://localhost:9870/dfshealth.html>



Started:	Fri Apr 17 22:44:51 +0300 2020
Version:	3.2.1, rb3cbbb467e22ea829b3808f4b7b01d07e0bf3842
Compiled:	Tue Sep 10 18:56:00 +0300 2019 by rohitsharmaks from branch-3.2.1
Cluster ID:	CID-11d9a063-fc7b-4208-b192-2e4b6b8736f
Block Pool ID:	BD-20320326446-168-1-168-1687462857460

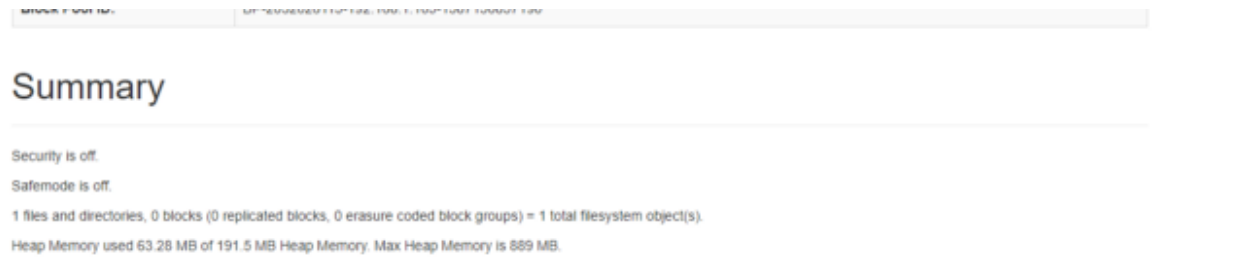


Figure 22 — Name node web page

- Data node web page: <http://localhost:9864/datanode.html>

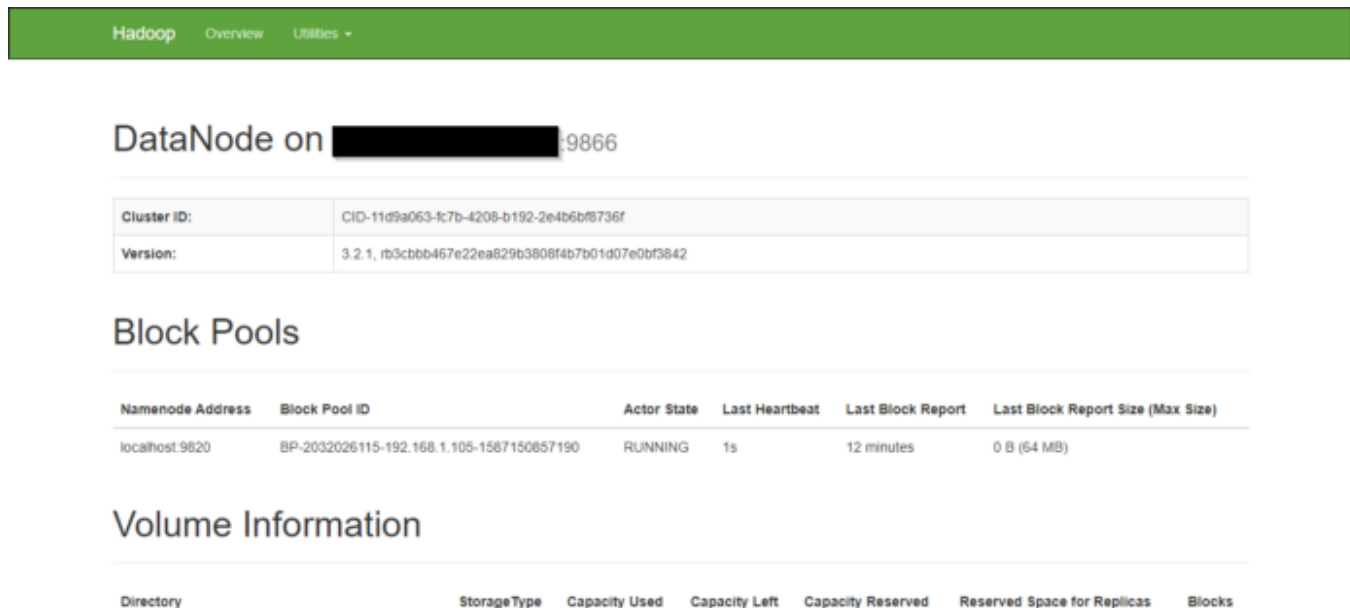


Figure 23 — Data node web page

- Yarn web page: <http://localhost:8088/cluster>

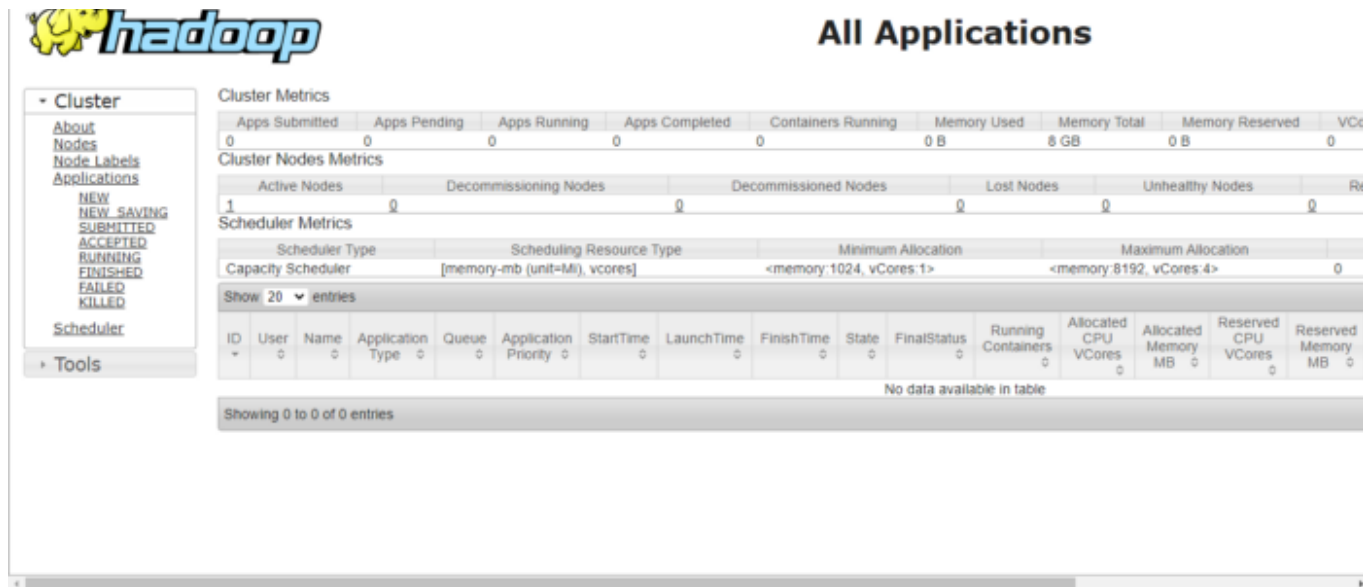


Figure 24 — Yarn web page

8. References

1. Hadi Fadlallah, [Installing Hadoop 3.1.0 multi-node cluster on Ubuntu 16.04 Step by Step](#), TowardsDataScience.com
2. Jozef Jarosciak, [How to install a Hadoop single node cluster on Windows 10](#)
3. Raymond Tang, [Install Hadoop 3.2.1 on Windows 10 Step by Step Guide](#), kontekst.tech
4. [Stack overflow Q/A website](#)

Sign up for The Daily Pick

By Towards Data Science

Hands-on real-world examples, research, tutorials, and cutting-edge techniques delivered Monday to Thursday. Make learning your daily ritual. [Take a look](#)

Your email

Get this newsletter

By signing up, you will create a Medium account if you don't already have one. Review our [Privacy Policy](#) for more information about our privacy practices.

Hadoop

Big Data

Windows 10

Hadoop Cluster

Hadoop 3

Learn more.

Medium is an open platform where 170 million readers come to find insightful and dynamic thinking. Here, expert and undiscovered voices alike dive into the heart of any topic and bring new ideas to the surface. [Learn more](#)

Make Medium yours.

Follow the writers, publications, and topics that matter to you, and you'll see them on your homepage and in your inbox. Explore

Share your thinking.

If you have a story to tell, knowledge to share, or a perspective to offer — welcome home. It's easy and free to post your thinking on any topic. [Write on Medium](#)

[About](#)

[Help](#)

[Legal](#)